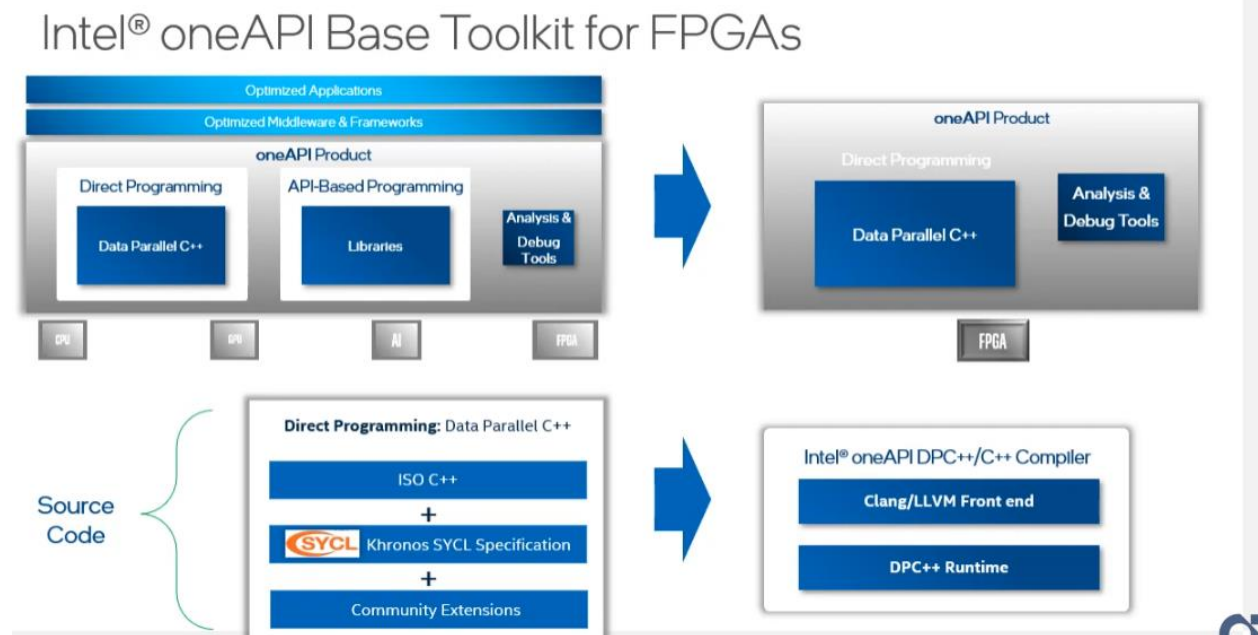
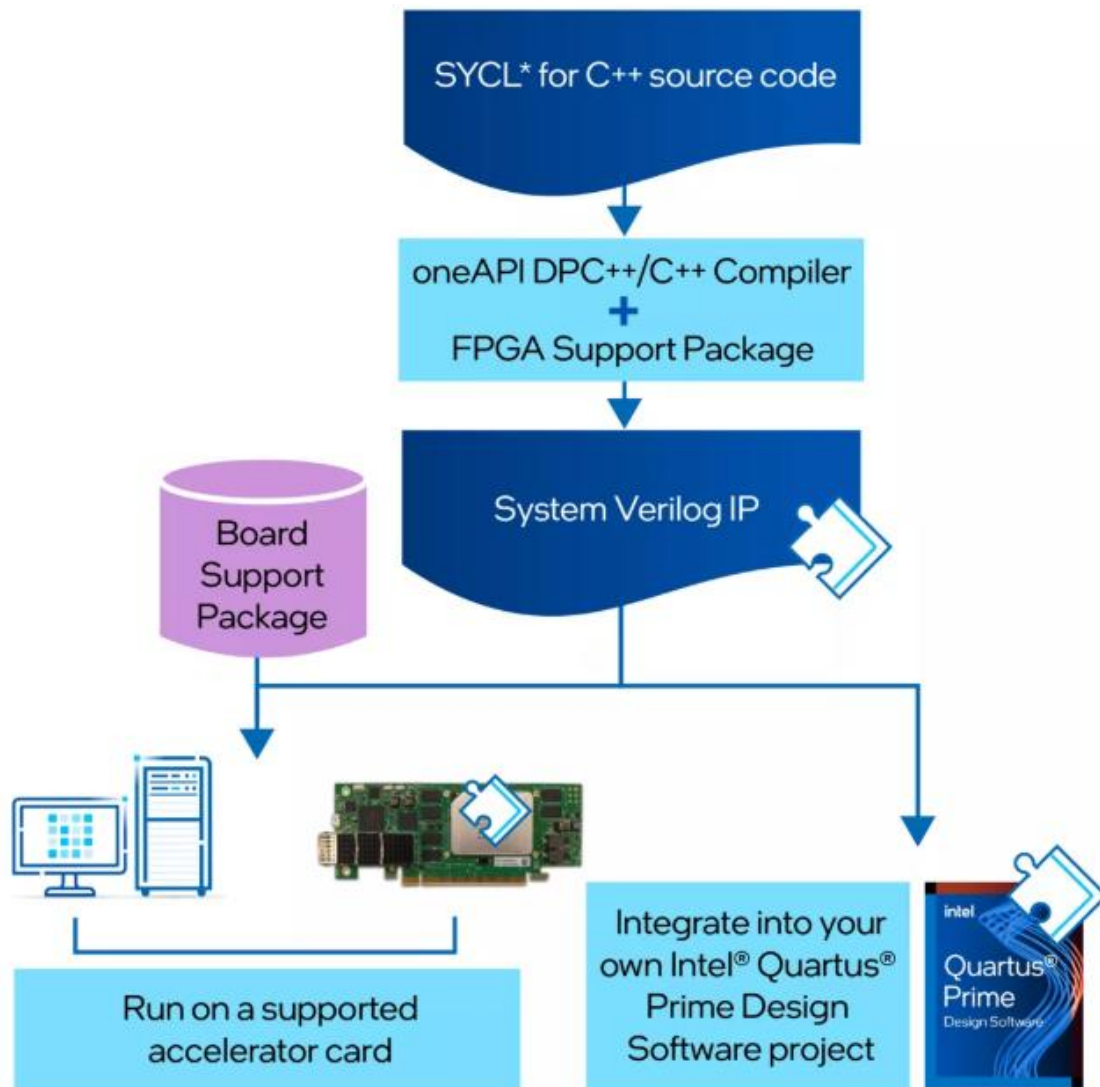


Here we are taking FPGA's, compiling and running form CPU is straight forward and GPU is very similar too because the accelerator device which is the CPU and GPU are available in the host device but the FPGA are a little bit different

To understand the general structure of DPC++ you can refer to the image below





Some general information about DPC++/Sycl:

- DPC++ is based on C++ and Sycl, and Sycl is based on OpenCL
- Even though it meant to target xPU's but you do still need to 'tune'

In writing DPC++/SYCL code there are three main scopes to consider

## DPC++/SYCL: Three Scopes

- DPC++ Programs consist of 3 scopes:
  - **Application scope** - Normal host code
  - **Command group scope** - Submitting data and commands that are for the accelerator
  - **Kernel scope** - Code executed on the accelerator
- The full capabilities of C++ are available at application and command group scope
- At kernel scope there are limitations in accepted C++
  - Most important is no recursive code
  - See SYCL specification for complete list

```
void dpcpp_code(int* a, int* b, int* c) {  
    //Set up an FPGA device selector  
    ext::intel::fpga_selector selector;  
    // Set up a DPC++ device queue  
    queue q(selector);  
    // Setup buffers for input and output vectors  
    buffer buf_a(a, range<1>(N));  
    buffer buf_b(b, range<1>(N));  
    buffer buf_c(c, range<1>(N));  
    //Submit Command group function object to the queue  
    q.submit([&](handler &h){  
        //Create device accessors to buffers  
        accessor a(buf_a, h, read_only);  
        accessor b(buf_b, h, read_only);  
        accessor c(buf_c, h, write_only);  
        //Dispatch the kernel  
        h.single_task<VectorAdd>([&]() {  
            for (int i = 0; i < kSize; i++) {  
                c[i] = a[i] + b[i];  
            }  
        });  
    });  
}
```

**Application Scope**

**Command Group Scope**

**Kernel Scope**

## The "Runtime"

- The DPC++/SYCL runtime is the program running in the background to control the execution and data passing needs of the heterogeneous compute execution
- It handles:
  - Kernel and host execution in an order imposed by data dependency needs (discussed later)
  - Passing data back and forth between the host and device
  - Querying the device
  - Etc.

## DPC++/SYCL Class: `device`

- The `device` class represents the accelerators in a oneAPI system
- The `device` class contains member functions for querying information about the device
- The function `get_info` gives information about the device:
  - Name, vendor, and version of the device
  - Width for built in types, clock frequency, cache width and sizes, online or offline

```
// Get all of the devices a system is capable of operating
std::vector<device> my_devices = device::get_devices();
// Grab the first device to print info out for
device my_device = my_devices[0];
// Print the name of the first device
std::cout << "Device: " << my_device.get_info<info::device::name>() << std::endl;
```



## DPC++/SYCL Class: `queue`

- A `queue` is a mechanism where work is submitted to a `device`
- A queue submits command groups to be executed by the SYCL runtime
- A `queue.submit()` is the beginning of the command scope
  - Groups of work to be executed by the SYCL runtime on an accelerator
- A queue maps to a single device

```
// Declare a queue to a device
queue q(selector);

// Submit things to the queue
q.submit([&](handler& h) {
    // COMMAND GROUP CODE
});
```

The handler is a class that contains all of the command group functions of SYCL

You can think of it as an abstraction of the runtime

This keeps us from having to type handler:: again and again in the command group scope

## DPC++/SYCL Class: `kernel`

- The `kernel` encapsulates code that will be run on the accelerator
- A `kernel` object is not explicitly constructed by the user
- It is constructed when a kernel dispatch function, such as `parallel_for()` or `single_task()` is called

```
q.submit([&](handler& h) {  
  
    // The "kernel" is everything after the kernel dispatch function  
    h.single_task<VectorAdd>([&]() {  
  
        // Everything inside here is "KERNEL SCOPE"  
        for (int i = 0; i < kSize; ++i) {  
            c[i] = a[i] + b[i];  
        }  
    });  
});
```

## Single Task Kernels

- `single_task()` kernels allow complex or lengthy datapaths to be built from custom hardware in FPGAs
- Useful to offload code with dependencies that are difficult to execute in a data parallel fashion
- Look like CPU code
  - Contain an outer loop to process all data
- Ideal for & recommended for FPGAs

```
for(int i=0; i < 1024; i++){  
    a[i] = b[i] + c[i];  
};
```

CPU  
Implementation

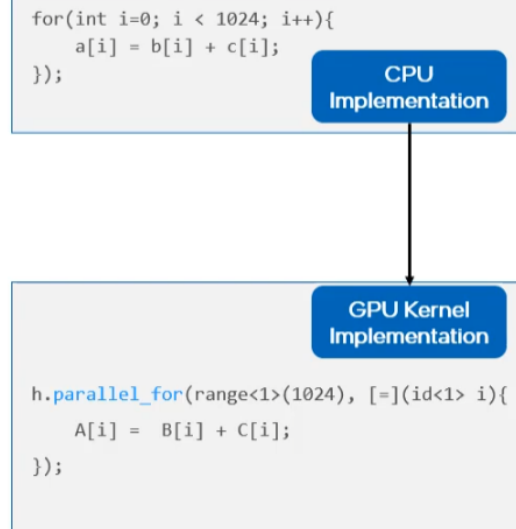


```
h.single_task([&]() {  
    for (int i=0; i < 1024; i++) {  
        A[i] = B[i] + C[i];  
    }  
});
```

FPGA Kernel  
Implementation

## Parallel Kernels

- Parallel Kernels allow multiple instances of an operation to execute in parallel
- Parallel kernels are expressed using the `parallel_for()` function
- Kernels that are easily expressed in this way do well on **GPUs**
  - Will be functional in an FPGA, but usually result in a less optimal implementation



## DPC++/SYCL Class: **buffer** and **accessor**

- **buffer**
  - Encapsulates data in a SYCL application
  - Across both devices and host!
- **accessor**
  - Mechanism to access buffer data
  - Determines data dependencies in that order kernel executions (covered later)

```
int main() {
    ... // Code to set up standard C++ vectors

    buffer buf_a(vector_a);
    buffer buf_b(vector_b);
    buffer buf_c(vector_c);

    queue q(selector);

    q.submit([&](handler& h) {
        accessor a(buf_a, h, read_only);
        accessor b(buf_b, h, read_only);
        accessor c(buf_c, h, write_only);

        h.single_task<VectorAdd>([=]() {
            for (int i = 0; i < kSize; i++) {
                c[i] = a[i] + b[i];
            }
        });
    });
};
```

You have to include “`sycl.hpp`” and “`fpga_extensions.hpp`”



# DPC++ / SYCL Simple Program Walk-Through

```
void dpcpp_code(int* a, int* b, int* c) {  
    //Set up an FPGA device selector  
    ext::intel::fpga_selector selector;  
    // Set up a DPC++ device queue  
    queue q(selector);  
  
    // Setup buffers for input and output vectors  
    buffer buf_a(a, range<1>(N));  
    buffer buf_b(b, range<1>(N));  
    buffer buf_c(c, range<1>(N));  
  
    //Submit Command group function object to the queue  
    q.submit([&](handler &h){  
        //Create device accessors to buffers  
        accessor a(buf_a, h, read_only);  
        accessor b(buf_b, h, read_only);  
        accessor c(buf_c, h, write_only);  
  
        //Dispatch the kernel  
        h.single_task<VectorAdd>([&]() {  
            for (int i = 0; i < kSize; i++) {  
                c[i] = a[i] + b[i];  
            }  
        });  
    });  
}
```

**Step 1:** Create a device selector targeting the FPGA

**Step 2:** Create a device queue, using the FPGA device selector

**Step 3:** Create buffers

**Step 4:** Submit a command for execution

**Step 5:** Create buffer accessors so the FPGA can access the data

**Step 6:** Send a kernel for execution


Done!

The contents of buf\_c are copied to \*c when the function finishes

(because of the buffer destruction of buf\_c)

```
int main() {  
    auto R = range<1>{ num };  
    buffer<int> A{ R }, B{ R };  
    queue Q;  
  
    Q.submit([&](handler& h) {  
        auto out = A.get_access<access::mode::read_write>(h);  
        h.parallel_for(R, [=](id<1> idx) {  
            out[idx] = idx[0]; });  
    }); Kernel 1  
  
    Q.submit([&](handler& h) {  
        auto out = A.get_access<access::mode::read_write>(h);  
        h.parallel_for(R, [=](id<1> idx) {  
            out[idx] = idx[0]; });  
    }); Kernel 2  
  
    Q.submit([&](handler& h) {  
        auto out = B.get_access<access::mode::read_write>(h);  
        h.parallel_for(R, [=](id<1> idx) {  
            out[idx] = idx[0]; });  
    }); Kernel 3  
  
    Q.submit([&](handler& h) {  
        auto in = A.get_access<access::mode::read>(h);  
        auto inout =  
            B.get_access<access::mode::read_write>(h);  
        h.parallel_for(R, [=](id<1> idx) {  
            inout[idx] *= in[idx]; });  
    }); Kernel 4  
}
```

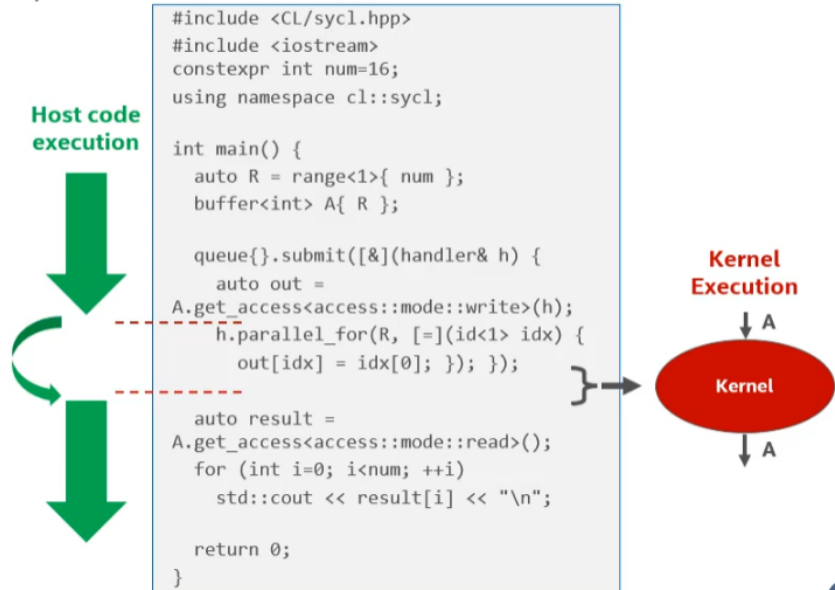
## Kernel Execution Order



- Kernels can execute at the same time
  - If no data dependencies
- Accessors are used to determine dependencies
- Execution ordering is automatically determined

# Asynchronous Host/Kernel Execution

- The execution of the host code is asynchronous to what is being executed on the accelerator
- If you need synchronization, you must impose that yourself

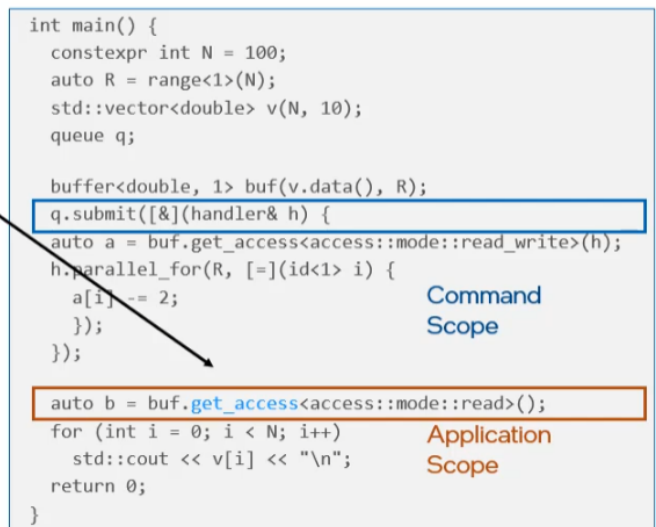


There are multiple ways to synchronize between the host and the kernel

You can make a dummy dependency in the application scope

## Synchronization Method 1: Host Accessor

- In the **command scope**,  
accessors are created for the accelerator
- In the **application scope**,  
accessors are created for the host
- A host accessor creates a dependency node in the execution graph
  - Execution at the host is blocked until the data is ready



Or you can depend on the buffer destruction because it has to wait until the kernel finishes execution to destruct the buffer



## Synchronization Method 2: Buffer Destruction

- Buffer creation happens within a separate function scope
- When execution advances beyond this function scope, buffer destructor is invoked
- Relinquishes ownership of data and copies back the data to the host memory
- Scope can also be created with simple use of `{ }`

```
#include <sycl/sycl.hpp>
constexpr int N=100;
using namespace cl::sycl;

void dpcpp_code(std::vector<double> &v, queue &q){
    auto R = range<1>(N);
    buffer<double, 1> buf(v.data(), R);
    q.submit([&](handler &h) {
        auto a = buf.get_access<access::mode::read_write>(h);
        h.parallel_for(R, [=](id<1> i) {
            a[i] -= 2;
        });
    });
}

int main() {
    std::vector<double> v(N, 10);
    queue q;
    dpcpp_code(v,q);
    for (int i = 0; i < N; i++)
        std::cout << v[i] << "\n";
    return 0;
}
```

Or you can do it using a different syntax, wrapping it in an inner scope using curly braces

Or simply use the `q.wait()`;

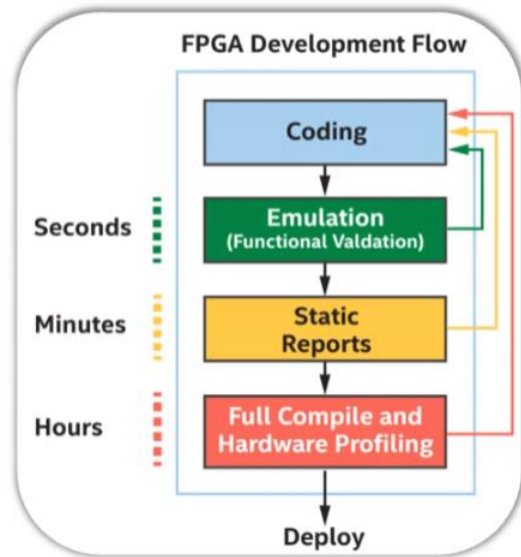
```
q.submit([&](handler &h) {
    // Create an accessor for each buffer with access permission: read, write or
    // read/write. The accessor is a mean to access the memory in the buffer.
    accessor a(a_buf, h, read_only);
    accessor b(b_buf, h, read_only);

    // The sum_accessor is used to store (with write permission) the sum data.
    accessor sum(sum_buf, h, write_only, no_init);

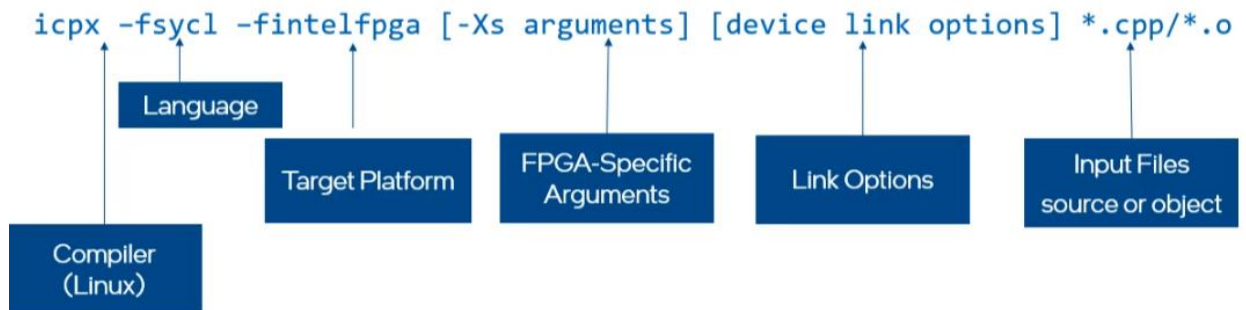
    // Use parallel_for to run vector addition in parallel on device. This
    // executes the kernel.
    // 1st parameter is the number of work items.
    // 2nd parameter is the kernel, a lambda that specifies what to do per
    // work item. The parameter of the lambda is the work item id.
    // SYCL supports unnamed lambda kernel by default.
    h.parallel_for(num_items, [=](auto i) { sum[i] = a[i] + b[i]; });
});
// Wait until compute tasks on GPU done
q.wait();
```

## FPGA Development Flow for oneAPI Projects

- FPGA Emulator target (Emulation)
  - Compiles in seconds
  - Runs completely on the host
- Optimization report generation
  - Compiles in seconds to minutes
  - Identify bottlenecks
- FPGA bitstream compilation
  - Compiles in hours
  - Enable profiler to get runtime analysis



## Anatomy of the SYCL\* Command Targeting FPGAs



## Emulation Command

```
#ifdef FPGA_EMULATOR
    ext::intel::fpga_emulator_selector device_selector;
#else
    ext::intel::fpga_selector device_selector;
#endif
```

Include this construct in  
your code

```
icpx -fyscl -fintelfpga -DFPGA_EMULATOR <source_file>.cpp
```



## Command to Produce an Optimization Report

```
icpx -fyscl -fintelfpga -Xhardware -fsycl-link<source_file>.cpp
```

The default value for `-fsycl-link` is `-fsycl-link=early`  
which produces an early image object file and report

- A report showing optimization, area, and architectural information will be produced in `<file_name>.prj/reports/`

## Bitstream Compilation



### Runs Intel® Quartus® Prime Software “under the hood” (no licensing required)

Developers can:

- Compile FPGA bitstream for their design and run it on an FPGA.
- Attain automated timing closure.
- Obtain In-hardware verification.
- Take advantage of Intel® VTune™ Profiler for real-time analysis of design.

## Compile to FPGA Executable with Profiler

```
icpx -fsycl -fintel FPGA -Xhardware <source_file>.cpp -Xsprofile
```

Optional:

If included, the profiler will be instrumented within the image and you will be able to examine data with the Intel® VTune™ Profiler when running the executable.

To compile to FPGA executable without profiling hardware, leave off `-Xsprofile`.

## What If Only the Host Code Changes?

- In the default case, the DPC++ Compiler handles generating the host executable, device image, and final executable
- What if you only change code that is run on the host, but nothing that affects the FPGA compile?

```
icpx -fsycl <source_file>.cpp -reuse-exe -Xshardware -fintelfpga
```

Normally, a compile to a full bit-stream takes hours. This flag instructs the compiler to re-use the FPGA bitstream if no code changes affect it.


A compile that would take hours now takes minutes...

# Optimizations

- Code to Hardware: An Introduction
- Loop Optimization
- Memory Optimization
- Reports
- Other Optimization Techniques

How Is a Pipeline Built?

- Hardware is added for
  - Computation
  - Memory Loads and Stores
  - Control and scheduling
    - Loops & Conditionals



```
for (int i=0; i<LIMIT; i++) {  
    c[i] = a[i] + b[i];  
}
```